

Recipe determinants of quantization fragility under image corruption: paired evidence and a regression-head localization in object detectors

Dinh Thuan Nguyen^a, Lam Phuong Nguyen^a, Vinh Huy Nguyen^b, Sy Vu Quang^c, Mohan Rajesh Elara^d and Anh Vu Le^{c,*}

^aFaculty of Electrical and Electronics Engineering, Ton Duc Thang University, Ho Chi Minh City, Vietnam

^bCenter of Visualization and Simulation, Duy Tan University, Da Nang, Vietnam

^cAdvanced Intelligent Technology Research Group, Faculty of Electrical and Electronics Engineering, Ton Duc Thang University, 19 Nguyen Huu Tho Street, Tan Phong Ward, District 7, Ho Chi Minh City 700000, Vietnam

^dROAR Lab, Engineering Product Development, Singapore University of Technology and Design, Singapore

ARTICLE INFO

Keywords:

Object detection
Post-training quantization
Robustness evaluation
Image corruption
Paired evaluation
Selective precision

ABSTRACT

Whether an 8-bit quantized detector stays reliable under corruption cannot be read off the format label: it depends on how the engine was built and which model region is quantized. We pair every corrupted input with a codec-matched clean control, so identical bytes reach executable TensorRT INT8 and FP8 engines, and estimate the corruption-associated change in the FP8–INT8 average-precision (AP) gap, ΔE , inside common-image bootstrap samples. Across nine dataset×scale blocks of YOLO11 on KITTI, VOC, and a 2,000-image COCO subset, ΔE is strongly heterogeneous ($I^2 = 87\%$) and blocks of both signs remain significant after Holm correction, so the corruption-induced change in the gap has no stable sign—even though FP8 exceeds INT8 on absolute clean and corrupted AP in every matched block mean. In RetinaNet, a matched-recipe decomposition localizes a severe INT8 collapse (−21.2 AP on KITTI) to the regression head: leaving only that head in floating point restores clean-competitive accuracy without retraining (49.60 vs. 49.48 AP for matched FP8 on original clean inputs), matching FP8 on the KITTI corrupt mean but leaving a residual that widens under corruption to about one AP point on VOC. A two-fold intervention, each fold frozen before its own runs, then calibrates INT8 on corrupted images and tests unseen families. Its frozen held-out criterion is met in two of the four fold×dataset cells, is insignificant in one, and fails with a significant across-the-board harm in the other, while a corruption-specific held-out benefit is significantly positive in exactly one fold×dataset cell (on both intervention tiers); codec-matched calibration controls (post hoc for fold 1, pre-specified for fold 2) bound the codec level contribution far below the full-INT8 level effects, though codec interaction effects are the same order as the corruption-specific components and are reported alongside them. Calibration coverage is therefore a context-dependent secondary factor, not a reliable repair for the corruption fragility, for which regression-head exclusion is the accuracy-validated fix (deployment cost not evaluated)—with the caveat that its strongest evidence comes from the development-contaminated KITTI check. Quantization robustness should be reported per recipe and per model region.

1. Introduction

Object detectors deployed on edge accelerators increasingly rely on post-training quantization (PTQ) to meet memory, energy, and latency budgets [1–5]. Two 8-bit formats dominate this setting: INT8, with a mature integer-arithmetic inference path [4, 6, 7], and FP8, which trades mantissa precision for exponent range on modern hardware [8, 9]. At the same time, deployment images deviate from benchmark conditions through noise, blur, weather, and compression artifacts that measurably degrade detector accuracy [10–13]. A natural question follows: which quantized format tolerates corruption better? We argue that this question, as usually asked, is ill-posed, because “INT8” and “FP8” name datatypes, whereas an executable detector is determined by a *recipe*: calibration data and preprocessing, quantizer placement and operator coverage, excluded regions, graph transformations, and runtime settings.

Two confounds make naive format comparisons misleading. First, a gap measured only on corrupted images

mixes the accuracy difference already present on clean images with the change induced by corruption; when one format also wins on clean images, the corrupted-input gap overstates the corruption effect. Second, recipes can differ systematically across formats in published pipelines: INT8 engines may be built with calibration distributions, preprocessing contracts, and operator-exclusion rules that differ from their FP8 counterparts—our own historical engines are one such instance (Section 4.2), and the two formats’ different numerical structure invites different recipe choices [14]—so a “format effect” may in fact be a recipe effect. A credible comparison must therefore pair treatments on identical image bytes, subtract the matched clean difference, and—crucially—vary the recipe deliberately rather than only the datatype.

This paper provides such a study on executable TensorRT engines evaluated end-to-end, with the decoder, non-maximum suppression (NMS), and evaluator held fixed within each detector family. We define a paired estimand ΔE : the change in the FP8–INT8 AP gap between a JPEG quality 95 clean control and each corruption cell, estimated inside common-image bootstrap resamples so

E-mail addresses: nguyendinhthuan.st@tdtu.edu.vn (D.T. Nguyen), nguyenvinhthuan.st@tdtu.edu.vn (L.P. Nguyen), nguyenvinhthuan@tdtu.edu.vn (V.H. Nguyen), vuquangsy@tdtu.edu.vn (S.V. Quang), rajeshelara@sutd.edu.sg (M.R. Elara), and leanhvu@tdtu.edu.vn (A.V. Le).

that pairing covariance is preserved. Under a frozen analysis contract—the estimands, condition schedule, and bootstrap plan are versioned before the paired analysis, while the YOLO11 engines themselves predate it and are disclosed as such (Section 3.4)—covering YOLO11n/m/x [15] on KITTI, VOC, and a deterministic 2,000-image COCO subset with 14 conditions per block, we find that ΔE reverses sign across dataset $\times$ model combinations (Section 4.1). The format label alone does not determine corruption behavior.

The sharpest evidence comes from a second detector family. On RetinaNet (ResNet-50 feature-pyramid-network (FPN) backbone) [16, 17], the historical INT8 recipe collapses on KITTI (−42.3 AP relative to a matched FP8 engine); repairing its calibration preprocessing contract still leaves a −21.2-AP deficit for matched INT8, and a stage-wise recipe decomposition traces that residual to one model region: quantizing only the classification head is harmless, quantizing only the regression head reproduces the collapse, and excluding the regression head alone restores clean-competitive INT8 accuracy (49.60 AP, matching FP8-matched at 49.48) without retraining (Section 4.2). The matched collapse is much milder on VOC (−2.9 AP), so the fragile region’s contribution is itself dataset-dependent.

Finally, we turn the diagnostic into a controlled two-fold intervention. If the INT8 corruption fragility were caused by calibration on clean images only—ranges that fail to cover corrupted-input activations—then calibrating on corruption-augmented images should restore accuracy on held-out corruption types. Freezing complementary splits (Gaussian noise and JPEG calibrated with fog and motion blur held out, and vice versa) shows that the intervention itself is context-dependent: the frozen held-out criterion is met on KITTI fold 1 (+1.07 AP) and VOC fold 2 but fails on the other two cells, including a significant across-the-board harm on KITTI fold 2, and a corruption-specific held-out component is significantly positive in exactly one cell. Corruption-aware calibration is therefore not supported as a reliable repair for the fragility, and clean-only calibration ranges are not supported as its dominant repairable cause; the regression head remains the dominant fragile component in this recipe, and its exclusion is the accuracy-validated retraining-free fix—validated on accuracy, not on deployment cost (Section 4.3).

This study makes five contributions:

- A paired, codec-controlled evaluation protocol for executable quantized detectors whose estimand isolates the corruption-associated change in a between-treatment gap, with pairing preserved by common-image bootstrap samples.
- Evidence, with a formal heterogeneity test and multiplicity control, that the corruption interaction is context-dependent within one detector family: ΔE changes sign across dataset and model-scale blocks, so the corruption-induced change in the gap has no stable sign on this protocol, while FP8 exceeds INT8 on absolute AP in every matched block mean.

- A recipe decomposition that localizes the RetinaNet INT8 collapse to regression-head quantization.
- A retraining-free selective-precision INT8 recipe—only the regression head kept in floating point—whose exclusion direction is validated under the frozen analysis contract on both datasets, while the >20-AP magnitude is a KITTI estimate carrying a flagged development-contamination caveat and a significant VOC residual is reported (deployment cost unmeasured).
- A controlled corruption-aware calibration intervention over two complementary in-family/held-out folds with codec-matched controls, analyzed with the same paired estimand: the intervention’s held-out effect is itself sign-unstable across fold $\times$ dataset cells, so calibration coverage contributes context-dependently and is not supported as a reliable repair for the corruption fragility.

Section 2 reviews quantization under corruption, detector-specific PTQ, and evaluation methodology. Section 3 defines the protocol, treatments, datasets, and the frozen intervention design. Section 4 reports the interaction landscape, the RetinaNet decomposition, and the intervention outcome. Section 5 discusses mechanism, practical guidance, and limitations. Section 6 concludes.

2. Related work

2.1. Post-training quantization for detectors

PTQ estimates numerical ranges with representative data or data-free surrogates and reduces quantization error through clipping, rounding, or local reconstruction [1, 18–23]. Detector-specific PTQ is harder than classification because classification confidence, box localization, feature pyramids, and foreground–background imbalance [24] must remain jointly calibrated: PD-Quant uses prediction-difference losses, DetPTQ a detection-output loss, and Reg-PTQ specializes calibration and quantization for regression structures [25–27]. InlierQ suppresses task-irrelevant activation anomalies using a small unlabeled calibration set [28]. These methods optimize clean-input accuracy; our decomposition instead asks which structural component of a deployed INT8 detector is responsible for fragility, a question Reg-PTQ’s emphasis on regression structures anticipates but does not measure at the executable-engine level.

Quantization-aware training and distillation approaches [29–32] can produce robust low-precision detectors but require retraining. Our selective-precision arm is deliberately retraining-free: it changes which regions of the graph are quantized, keeping the deployed checkpoint and calibration budget fixed, which makes the recipe contrast interpretable.

2.2. INT8, FP8, and executable treatments

Integer-only inference is an established path [4, 6, 33], while FP8 formats evolved from training formats [34, 35] to practical PTQ layouts [8, 9, 36, 37]. Comparisons show that accuracy and hardware efficiency can favor different

representations and that the preferred 8-bit format varies by layer and architecture [14, 38]. The nominal datatype therefore does not describe an executable engine: quantizer coverage, scale resolution, excluded operators, calibration, and runtime all contribute. Our protocol records the complete ModelOpt/TensorRT treatment and manipulates recipe components one at a time, which is what allows the recipe-format confound to be measured rather than assumed.

2.3. Corruption robustness and quantized models

Corruption benchmarks [10] measure aggregate degradation for classifiers, and detector studies report task-specific failure under weather, noise, blur, and compression [11–13]. The standard summary normalizes mean performance under corruption (mPC) by clean performance (P_{clean}) as the relative rate $\text{rPC} = \text{mPC}/P_{\text{clean}}$ [11], which answers a within-model retention question. Our estimand instead asks how the *between-treatment* AP-point gap changes on identical corrupted bytes; a related absolute-AP guard prevents a shrinking gap near a shared failure floor from being read as robustness.

Work linking quantization and distribution shift is comparatively thin: Karimov et al. evaluate quantized detectors under degraded inputs and report sensitivity to the evaluation setting [39]; a separate line of work measures quantization’s effect on *adversarial* robustness [40], a distinct shift family. TTAQ adapts quantization during continual shift [41], Recti-Q corrects out-of-distribution quantized features [42], and test-time adaptation more broadly updates the model at inference, including normalization-statistics adaptation under corruption [43–46]. Corruption-aware training and augmentation improve robustness by exposing optimization to degraded inputs [47, 48], and detector error analyses show that conclusions depend on how errors are decomposed [49, 50]; evaluation-protocol work likewise shows that evaluation design itself changes system-level conclusions [51, 52]. Our intervention is related in spirit to corruption-aware calibration, but differs in the experimental control: the calibration corruption families are disjoint from the held-out evaluation families, so any claimed robustness benefit must generalize rather than memorize.




Figure 1: Paired evaluation protocol. A deterministic clean control encoded at JPEG quality 95 (J95) and each corrupted input are created before the precision branch. The four paired cells combine INT8 and FP8 treatments with J95-clean and corrupted inputs; the FP32 engine is a separate reference and cancels from the interaction ΔE . Common image resamples preserve pairing. Clean-control fidelity and absolute corrupted AP—and the signed loss $D = AP_{J95} - AP_c$ shown in the floor guard—are checked before an interaction is interpreted.

3. Methodology

3.1. Executable treatments as recipes

Let a *treatment* denote the complete recorded configuration of an executable detector: checkpoint, Open Neural Network Exchange (ONNX) graph, quantizer placement, quantize/dequantize (Q/DQ) coverage, calibration list, calibration preprocessing, excluded operators, TensorRT builder settings, runtime, and decoder. Throughout, “INT8” and “FP8” qualify such treatments, not isolated datatypes, and an *arm* denotes one treatment instance within an experiment. All engines were built with TensorRT 11.1.0.106 [53] on an NVIDIA GeForce RTX 5090 with a 4,096 MiB builder workspace, from ONNX graphs quantized by NVIDIA ModelOpt using entropy calibration [54, 55]; the study-built registries record ModelOpt 0.45.0, while the pre-study YOLO11 ONNX registries predate that version field. Builder policy differs across the two evidence programs and is itself a recipe attribute, recorded per engine in the manifests: the YOLO11 engines predate the frozen contract and were built by `trtexec` with TensorRT defaults (TF32 enabled), while the RetinaNet quantization arms built for this study use the TensorRT Python builder with TF32 disabled (`tf32: false` in each engine registry). The RetinaNet FP32 reference and the FP8/INT8-legacy arms are pre-study `trtexec` builds that carry TensorRT defaults, so TF32 is enabled there; within each family the TF32 state is identical across the paired *quantized* arms, which are the only arms entering the ΔE contrasts. Calibration lists contain 512 deterministically

selected training images for all paired and intervention arms (the earlier diagnostic sweeps include 128-image and max-estimator variants, marked in Table 4); evaluation and calibration images are disjoint.

Within this framing we define the RetinaNet recipe arms used below:

- **INT8-legacy:** the historical transferred recipe; its calibration preprocessing does not match inference preprocessing (a contract violation retained as a diagnostic arm).
- **INT8-matched:** entropy INT8 with calibration preprocessing identical to inference (`retinanet_normalized`) and the same 512-image budget as FP8.
- **INT8-selective:** INT8-matched with all `/head/regression_head/` nodes left in floating point—a selective-precision recipe requiring no retraining.
- **INT8-corruptcalib** and **INT8-sel-corruptcalib:** the matched and selective recipes rebuilt on a corruption-aware calibration set (noise and JPEG; Section 3.6); **INT8-cc2calib** and **INT8-sel-cc2calib** are the complementary fold (fog and motion blur).
- **INT8-q95calib** and **INT8-sel-q95calib:** codec-matched calibration controls—the same 512 clean images re-encoded at JPEG quality 95 with no corruption (Section 3.6).
- **FP8-matched** and **FP32:** matched-recipe FP8 and the high-precision reference.

Two pre-study engines are cited as context only and are outside this arm list: INT8-legacy above and the **FP8-legacy**

exploratory build (a `trtexec`-default, TF32-on engine evaluated under the frozen protocol for reference). Three senses of “matched” recur and are kept distinct: *contract-matched* arms use calibration preprocessing identical to inference; *codec-matched* inputs share the terminal JPEG quality-95 encoding; and *paired* comparisons evaluate treatments on identical images (and identical resamples in the bootstrap).

3.2. Paired estimands and interpretation order

For dataset d , model m , quantized treatment q , and condition c , write $AP_{d,m,q,c}$ for COCO-style $AP@.5:.95$ reported in AP points. J95 denotes the clean control: the source image deterministically encoded at JPEG quality 95 with subsampling disabled and no named corruption (Fig. 1). The between-treatment gaps on clean and corrupted inputs are

$$G_{d,m,x} = AP_{d,m,FP8,x} - AP_{d,m,INT8,x}, \quad x \in \{J95, c\}, \quad (1)$$

and the primary interaction is

$$\Delta E_{d,m,c} = G_{d,m,c} - G_{d,m,J95}. \quad (2)$$

Positive ΔE means the FP8–INT8 gap widens under corruption relative to the clean baseline; negative means it narrows. The FP32 reference cancels algebraically. We abbreviate $G_0 = G_{d,m,J95}$ and write G_c (“corrupt-mean gap”) for the equal-weighted mean of $G_{d,m,c}$ over the 12 corruption–severity cells; the reported per-block ΔE is $G_c - G_0$ (Eq. (2)). Absolute corrupted AP is reported beside ΔE ; a shrinking gap near a shared accuracy floor is a floor effect, not robustness (the “floor guard”: interpretations are made jointly with the corrupted AP level). The contrast is descriptive, not causal—algebra analogous to a difference-in-differences interaction [56] without causal identification claims.

The same construction applies to recipe contrasts: we substitute two arms of the same format (e.g., INT8-selective minus INT8-matched) to isolate a recipe component under an identical protocol. This is the key methodological step that turns a format comparison into a recipe experiment.

3.3. Corruption suite and codec control

Four synthetic corruptions are generated deterministically: Gaussian noise ($\sigma \in \{8, 20, 40\}$ on the 0–255 scale), linear motion blur (kernel lengths $\{5, 13, 23\}$ at a per-image uniform angle), synthetic fog (strength–radius pairs $(0.18, 6)$, $(0.38, 12)$, $(0.62, 20)$ blended over a per-image random field), and JPEG compression (qualities $\{70, 40, 15\}$). Severity labels 1/3/5 order settings within each type. Every output is re-encoded at JPEG quality 95, so the clean control and corrupted inputs share the terminal codec; a SHA-256-derived seed binds dataset, image identifier, corruption, and severity. Within each detector family, inference fixes the input resolution (640-pixel letterbox), confidence threshold, decoder, and non-maximum suppression (NMS): YOLO11 uses a confidence threshold of 0.001 and class-wise NMS at intersection-over-union (IoU) 0.7 with at most 300 detections; RetinaNet uses 0.05 and batched NMS at IoU 0.5, also capped at 300 detections. AP evaluation uses the COCO protocol (`pycocotools`) with `maxDets=100`. Decoder

constants therefore differ across families, and cross-family comparisons compare recipe bundles rather than a shared decoder; all within-family contrasts share identical decoder settings.

3.4. Datasets and the frozen paired protocol

Three evaluation universes anchor the frozen paired analysis: KITTI validation (1,496 images) [57], PASCAL VOC2007 test (4,952 images) [58], and a deterministic 2,000-image subset of COCO val2017 [59] (seed 20260924; the subset is identical across all arms so that COCO comparisons remain paired). VOC and KITTI annotations were converted once to COCO format and frozen; reported values are COCO-style measurements on converted annotations. One checkpoint per dataset×model cell is used throughout, and its provenance is disclosed rather than idealized: the KITTI and VOC YOLO11 scales were fine-tuned from Ultralytics pretrained weights under the pilot program (runs `*_train_v1`), with the best checkpoint selected on the same validation partition that serves as the evaluation universe here—so KITTI/VOC absolute AP levels carry selection optimism whose residual effect on paired contrasts is unquantified; COCO YOLO11 uses the official pretrained checkpoint with no dataset fine-tune; the RetinaNet ResNet-50-FPN-v2 model was fine-tuned per dataset for 100 epochs at 640-pixel input (seed 20260807) with best-checkpoint selection by minimum validation loss on the same evaluation partition, under the same caveat. All checkpoint, split, and training-run manifests are hash-bound (Supplement S4). Each dataset×model block evaluates clean, codec-control, and the 12 corruption–severity cells on identical image ordering: 14 conditions per arm, except the RetinaNet FP32 and legacy arms, which predate codec control and carry 13. RetinaNet arms add the intervention engines on the same protocol (Section 3.6).

Two scope notes qualify interpretation. First, the KITTI validation universe was used while developing the selective-precision recipe, so KITTI selective-arm numbers are development-contaminated; VOC is the cleaner cross-dataset check for that arm. Second, coverage is asymmetric across families: RetinaNet runs on KITTI and VOC, YOLO11 on all three universes, so COCO evidence is single-family and cross-family comparisons on KITTI/VOC share the same checkpoints-selection caveat.

3.5. Paired image bootstrap

Uncertainty comes from the finite evaluation images with checkpoints, engines, calibration, decoder, and corruptions fixed; evaluation-set composition alone is a known source of score variance [60, 61]. Within each dataset×model block we use one deterministic schedule of 2,000 bootstrap resamples of image indices [62] shared across all arms and all conditions, so every paired contrast is evaluated inside common-image samples and the induced covariance is preserved. Reported point estimates are full-sample plug-in values; reported intervals are 2.5th–97.5th percentiles of the bootstrap distribution of each contrast, marked * when they exclude zero. Two-sided bootstrap

p -values, $p = \min\{1, 2 \min[\Pr(d \leq 0), \Pr(d \geq 0)]\}$ (smallest nonzero value $1/1000$ at $B = 2000$), are Holm-adjusted across the nine paired YOLO11 blocks [63]. Heterogeneity of ΔE across blocks is tested with Cochran’s Q on inverse-variance weights from the bootstrap standard errors [64, 65]; because blocks within a dataset share images, their estimates may be correlated; the cross-block covariance is not estimated, so Q is an independence approximation whose error direction is not established (shared-image resampling plausibly induces positive covariance, which would make Q conservative, but this is not verified). Diagnostic decomposition arms are single-point estimates without intervals; the exploratory layer’s cells carry their own bootstrap intervals in the retained ledger, as summarized in Section 3.7.

3.6. Corruption-aware calibration intervention

The intervention tests one hypothesis under decision rules frozen in a versioned configuration before any intervention run: *if the INT8-matched arm’s corrupted-input accuracy is limited by calibration images that never exhibit corrupted-input activation ranges, then corrupting the calibration distribution should recover held-out corruption AP*. This hypothesis is deliberately scoped to corruption robustness: it cannot by itself account for the clean-image collapse, which the decomposition attributes elsewhere. Each of the 512 calibration images is deterministically assigned—by a hash of dataset and image path under schedule `nn-corruptcalib-v1`—to exactly one of six in-family cells $\{\text{Gaussian noise, JPEG}\} \times \{1, 3, 5\}$ (77–102 images per cell) and corrupted with the same frozen parameters and JPEG quality-95 output encoding as the evaluation caches (Section 3.3). Fog and motion blur are never calibrated in this first fold; they form its held-out split. A complementary second fold, frozen in its own versioned configuration before any of its runs, swaps the roles: it calibrates the same 512 images on fog and motion blur (75–95 images per cell) and holds out Gaussian noise and JPEG, and it pre-specifies both decision rules for the full tier (INT8-cc2calib)—the fold-1 held-out criterion and the held-out ΔE rule against both INT8-matched and INT8-q95calib—while the fold-2 selective-tier contrasts are post-hoc. Its arms are INT8-cc2calib and INT8-sel-cc2calib. The fold-2 primary rule is worded non-directionally (“CI95 excluding 0 $\Rightarrow$ fold-1 frozen criterion replicated”); we read “replicated” through the directional fold-1 criterion it cites, so a significantly *negative* held-out change counts as failure. The fold-1 frozen configuration carries two success criteria: (i) full INT8 corrupt-calibration improves the *held-out* corrupt mean over INT8-matched with an interval excluding zero on the positive side, which would support the range-mismatch hypothesis; and (ii) the selective corrupt-calibration arm’s held-out mean *includes* zero relative to FP8-matched, which would let calibration explain the VOC residual gap. Criterion (ii) is deliberately null-seeking: an interval containing zero bounds the residual rather than proving equivalence, since no equivalence margin was

pre-specified. In-family improvement alone is evidence of calibration-fitting, not robustness transfer. Both criteria are reported regardless of outcome; where a criterion is stated as “excluding zero” we always read it directionally, so a significantly *negative* held-out change counts as failure, not success. Two interpretive analyses accompany the fold-1 criteria and are labeled as such because they were added after the fold-1 freeze (fold 2 subsequently froze the ΔE rule, but not the magnitude comparison): (i) a magnitude comparison of any held-out gain against the residual gap to matched FP8, and (ii) the paper’s own paired estimand applied to the calibration contrast, $\Delta E_{\text{held-out}} = (\text{held-out gain}) - (\text{J95 clean gain})$, which separates a corruption-specific benefit from a uniform accuracy shift that would appear on clean inputs too. This last subtraction rests on an additive, parallel-shift assumption—that any calibration-induced level shift on clean inputs transfers comparably to corrupted inputs—so it is an interpretive decomposition, not a causal identification (Supplement S7).

Two controls discipline interpretation. First, the corrupt-calibration images carry the same JPEG quality 95 terminal encoding as the evaluation caches, whereas the matched clean-calibration list points at original bytes; a gain of the intervention arm over INT8-matched is therefore potentially confounded by codec adaptation rather than corruption coverage. We isolate this component with the INT8-q95calib and INT8-sel-q95calib control arms, calibrated on the same 512 images re-encoded at JPEG quality 95 without corruption. These codec controls were added after the fold-1 freeze and after the fold-1 results were computed, when the confound was identified during review of the intervention design; they are therefore post-hoc controls on the fold-1 contrast and are labeled as such, while fold 2 pre-specifies its ΔE rule against the q95-matched anchor in advance. Second, the codec confound works in the direction that matters for the primary test: a confound that *inflates* the intervention arm’s gains biases the primary superiority test *toward* the positive verdict, so a positive held-out result cannot be attributed to corruption coverage until it survives against the codec-matched anchor. For the null-seeking selective criterion the direction of the confound reverses: codec adaptation that lifts the selective arm moves its negative held-out residual toward zero, so it could in principle manufacture an interval containing zero—a met criterion would be suspect, while an observed failure is conservative.

3.7. Evidence layers

Three primary evidence layers are kept distinct (Table 1): *diagnostic* evidence (Section 4.2) consists of single-point recipe sweeps that localized the fragile region during development; *frozen paired analysis* is the paired protocol of Section 3.4 with bootstrap intervals—the freeze applies to the evaluation and analysis contract (estimands, condition schedule, bootstrap schedule), not to the artifacts, since the YOLO11 engines, predictions, and evaluation partitions were built for the exploratory program and are shared with

Table 1

Evidence layers and their inferential status. Intervals exist only where pairing and the frozen schedule are defined.

Layer	Evidence	Engines/protocol	Uncertainty
Diagnostic	Region-wise recipe sweeps	Development builds	Single-point AP
Frozen paired analysis	Paired 14-cell protocol	Pre-existing engines under a frozen analysis contract	Paired bootstrap (2,000 resamples)
Intervention	Calibration interventions	Frozen fold configs; fold-1 code controls post hoc	Paired bootstrap (2,000 resamples)
Historical exploratory	144-cell prior grid	Shared engines; separate retrained-holdout and legacy builds	Reported as context

the historical layer; *intervention* evidence is the corrupt-calibration experiment and its codec controls (Section 3.6), whose two corruption-calibration folds were frozen prospectively before their engines were built while the fold-1 codec-control arms are post-hoc additions. A fourth, historical exploratory layer—the same paired YOLO11 engines and estimand analyzed earlier as a 144-cell grid over four datasets with a different scope and analysis plan—is reported only as context (Section 4.4).

4. Results

4.1. The FP8–INT8 interaction is not universal

Table 2 and Fig. 2 report the paired YOLO11 landscape: nine dataset×model blocks, each with the clean-control gap G_0 , the corrupt-mean gap G_c , and the paired interaction ΔE with its 95% interval. FP8 exceeds INT8 on clean-control AP in all nine blocks (G_0 from +0.68 to +3.97 AP points) and on the 12-cell corrupt mean as well (FP8 and INT8 corrupted AP columns, Table 2); the gap is positive on the level in every block, although a minority of individual severity cells reverse it. What lacks a stable sign is the change corruption induces in that gap: it is positive on VOC/YOLO11n (+0.31 [+0.07, +0.59]) and COCO-2k/YOLO11m (+0.60 [+0.27, +0.97]), negative on VOC/YOLO11m (−0.57 [−0.80, −0.34]), KITTI/YOLO11m (−0.68 [−1.32, −0.05]), KITTI/YOLO11x (−0.69 [−1.13, −0.33]), and COCO-2k/YOLO11x (−0.53 [−0.93, −0.18]), and compatible with zero on the remaining three blocks (unadjusted intervals; multiplicity is addressed below). Within VOC alone, the sign of the estimate flips with model scale (n : +0.31 excluding zero; m : −0.57 excluding zero; x : +0.14 null). A claim that INT8 is uniformly less robust under corruption—or uniformly comparable—is not supported on this protocol’s own evidence; within a fixed recipe, the sign of the interaction varies with the dataset×model context.

The heterogeneity is formal, not an artifact of reading nine intervals. Cochran’s test rejects a common ΔE across blocks ($Q = 63.3$ on 8 degrees of freedom, $p \approx 1 \times 10^{-10}$; $I^2 = 87\%$, between-block standard deviation $\tau = 0.44$ AP), so a pooled “format effect” would summarize nothing. The heterogeneity is not only between datasets: it holds within VOC ($Q = 29.3$, 2 df, $p \approx 4 \times 10^{-7}$) and within the COCO subset ($Q = 21.9$, 2 df, $p \approx 2 \times 10^{-5}$), where only model scale varies, though not within KITTI ($Q = 2.4$, $p = 0.30$). After Holm adjustment across the nine blocks, 4 blocks remain significant and they carry *both* signs—one positive (COCO-2k/YOLO11m) and three

negative (KITTI/YOLO11x, VOC/YOLO11m, COCO-2k/YOLO11x; Table 2); a DerSimonian–Laird pooled mean is −0.18 AP (−0.50, +0.13], a Wald interval), a level summary, not a common interaction. Within one dataset, VOC/YOLO11n and VOC/YOLO11m differ by +0.88 AP ($z = 4.9$ under an independence approximation that is conservative for positively correlated blocks). The heterogeneity also exists *within* blocks at the family level—on VOC/YOLO11x the JPEG family contributes −1.07 AP while Gaussian noise and motion blur contribute +0.69 and +0.67 AP—and excluding the saturated motion-blur-s5 cell leaves the cross-block heterogeneity essentially unchanged ($Q \approx 60.2$, $I^2 \approx 87\%$) and every Holm-surviving block’s sign intact, although the two KITTI blocks that do not survive Holm drift (YOLO11n’s null point estimate flips from −0.01 to +0.15 and YOLO11m’s unadjusted-significant negative turns null; Supplement S8, Table S2). The codec control also behaves as required: J95 clean AP tracks original clean AP within 0.3 points for every arm in these nine blocks, and the within-model retention statistic (corrupt mean over J95 clean) ranges 0.68–0.79 across blocks with INT8-versus-FP8 differences of at most ~ 0.02 —the pairing basis itself does not penalize either treatment.

4.2. A matched recipe exposes and localizes RetinaNet fragility

The non-YOLO comparator is where recipe effects dominate. Table 3 and Fig. 3 show the RetinaNet arms on KITTI and VOC under the frozen protocol. Three observations matter.

First, about half of the historical INT8 recipe’s collapse is a contract violation, not a format property: in the controlled diagnostic—which rebuilds the mismatched contract under the study’s common no-TF32 builder policy at the same 128-image calibration budget and changes only the preprocessing—KITTI J95 AP moves from 6.72 to 27.91, recovering 21.2 of the 41.8-point gap to the diagnostic FP8 build. On the frozen protocol’s original-clean inputs the historical engine reads 7.21 against the matched arm’s 28.29 (Table 3), a 42.3-point total deficit to matched FP8; that pair additionally differs in builder policy (the historical engine predates the study and carries TensorRT’s default TF32 state), so the diagnostic comparison is the attributional one. The legacy recipe must therefore be described as a mismatched-calibration artifact, not evidence about INT8 per se.

Second, the residual INT8-matched collapse is severe but dataset-dependent: −21.2 AP versus matched FP8 on KITTI but only −2.9 AP on VOC (original clean images, Table 3; −21.1 and −3.0 on the J95 basis). The paired interaction ΔE (FP8–INT8-matched, corrupt mean) is −4.69 [−5.49, −3.86] on KITTI and −0.26 [−0.58, +0.06] on VOC; the KITTI narrowing is consistent with a scale artifact of the additive estimand: INT8-matched enters corruption 21 points below FP8, and the two arms’ retention

Table 2

Frozen paired-analysis contrasts, FP8 minus INT8, per dataset×model block. FP8/INT8 corr. AP: absolute 12-cell corrupt-mean AP (the floor guard); G_0 : clean-control (J95) gap; G_c : corrupt-mean gap (the difference of the two absolute columns, computed before they are rounded to 0.1 AP); $\Delta E = G_c - G_0$ with paired-bootstrap interval. * interval excludes zero; p_{Holm} : bootstrap p -value Holm-adjusted across the nine blocks.

Dataset	Model	FP8 corr. AP	INT8 corr. AP	G_0	G_c	ΔE	95% interval	p_{Holm}
KITTI	YOLO11n	44.1	41.9	+2.14	+2.13	-0.01	[-0.67, +0.91]	0.914
KITTI	YOLO11m	50.4	47.1	+3.97	+3.29	-0.68*	[-1.32, -0.05]	0.124
KITTI	YOLO11x	51.0	50.3	+1.44	+0.74	-0.69*	[-1.13, -0.33]	0.008
VOC	YOLO11n	44.8	43.8	+0.68	+0.99	+0.31*	[+0.07, +0.59]	0.075
VOC	YOLO11m	51.8	50.7	+1.70	+1.12	-0.57*	[-0.80, -0.34]	<0.009
VOC	YOLO11x	54.3	53.3	+0.91	+1.05	+0.14	[-0.07, +0.36]	0.354
COCO-2k	YOLO11n	28.2	27.4	+1.10	+0.78	-0.32	[-0.67, +0.01]	0.168
COCO-2k	YOLO11m	38.0	36.5	+0.89	+1.49	+0.60*	[+0.27, +0.97]	0.008
COCO-2k	YOLO11x	42.4	41.9	+1.00	+0.47	-0.53*	[-0.93, -0.18]	0.018




Figure 2: Paired corruption interaction ΔE (FP8-INT8) with 95% common-image-bootstrap intervals. Solid markers denote intervals excluding zero (unadjusted); open diamonds mark blocks surviving Holm adjustment across the nine blocks ($p_{\text{Holm}} < 0.05$). The sign reverses across dataset×model blocks, so the interaction cannot be summarized as a single format-level effect.

rates are nearly equal (0.72 versus 0.75), so a similar relative decline narrows the absolute gap mechanically.

Third, the decomposition localizes the collapse to one model region (Table 4, diagnostic single-point arms on KITTI): quantizing only the classification head preserves accuracy ($50.57 \approx \text{FP32}$), quantizing only the regression head reproduces the collapse (16.41), and INT8 everywhere *except* the regression head recovers to 49.50—within 0.94 of the diagnostic FP32 build on the same J95 basis (50.44; the paired-protocol FP32 reference in Table 3, evaluated on original clean bytes rather than J95, is 50.15) and above full-graph FP8. Across all eleven INT8 diagnostic arms the outcome is bimodal in one variable: every arm that quantizes the regression head lands between 6.72 and 28.74 AP, whatever its calibration contract, estimator, budget, or site placement,

and every arm that leaves it in floating point lands between 49.14 and 50.57 AP. Two alternatives a reader might raise are inconsistent with this split. Position in the graph is not the explanation, because the classification head is equally terminal yet harmless when quantized alone. Quantization coverage is not the explanation either: the restricted Conv+Add placement (15.75) quantizes fewer sites than full INT8 and does worse, while backbone+FPN quantization with both heads in floating point (49.14) quantizes most of the network and does well. Effects are non-additive (quantizing both heads, 17.09, is worse than quantizing the full graph), so the arms identify a dominant fragile component rather than additive contributions. These arms are single engine builds from one calibration draw on the development set; the frozen paired

Table 3

RetinaNet arms under the frozen protocol (AP points, COCO-style AP on converted annotations). Corrupt mean = equal weight over the 12 corruption-severity cells. J95: codec-control clean; FP32 and legacy arms predate codec control (dashes). The INT8-legacy arm is the pre-study historical engine and additionally differs from the matched arms in builder policy (TensorRT-default TF32 state versus TF32 disabled in the TensorRT Python builder); the controlled preprocessing-only contrast is the diagnostic rebuild in Table 4.

Dataset	Arm	Clean AP	J95 AP	Corrupt mean AP
KITTI	FP32 reference	50.15	—	37.80
	FP8-matched	49.48	49.25	36.75
	INT8-legacy (contract mismatch)	7.21	—	7.43
	INT8-matched	28.29	28.18	20.37
	INT8-q95calib (codec control)	28.26	27.85	20.42
	INT8-selective (reg-head FP32)	49.60	49.53	36.58
	INT8-sel-q95calib	50.28	49.79	36.94
	INT8-corruptcalib	29.69	29.74	21.71
	INT8-sel-corruptcalib	49.55	49.03	36.08
	INT8-cc2calib (fold 2)	25.36	25.47	18.68
	INT8-sel-cc2calib (fold 2)	49.72	50.54	37.20
VOC	FP32 reference	57.23	—	42.32
	FP8-matched	56.17	56.21	41.48
	INT8-legacy (contract mismatch)	33.61	—	24.32
	INT8-matched	53.30	53.18	38.71
	INT8-q95calib (codec control)	53.33	53.39	38.54
	INT8-selective (reg-head FP32)	55.87	55.79	40.52
	INT8-sel-q95calib	55.81	55.95	40.40
	INT8-corruptcalib	53.78	53.91	39.35
	INT8-sel-corruptcalib	56.36	56.43	41.13
	INT8-cc2calib (fold 2)	54.15	54.22	39.97
	INT8-sel-cc2calib (fold 2)	56.40	56.43	41.58

Table 4

Diagnostic recipe decomposition on KITTI (single-point AP on the J95 codec-control inputs, 512-image calibration unless stated). These arms guided the paired-protocol design and carry no intervals; the “INT8 matched preprocessing, 512-img calibration” row reproduces the paired INT8-matched cell’s numbers from the same quantized ONNX graph (identical predictions; the serialized engine is a separate build, so the engine hash differs) rather than duplicating it.

Recipe arm (KITTI diagnostic, single-point)	J95 AP
INT8 legacy (mismatched calibration contract, 128-img)	6.72
INT8 matched preprocessing, 128-img calibration	27.91
INT8 matched preprocessing, 512-img calibration	28.18
INT8 max-estimator, 128-img calibration	23.94
INT8 restricted to Conv+Add sites, 128-img calibration	15.75
INT8 at exactly the FP8 compute sites (shared mask)	28.74
INT8 backbone+FPN only (heads FP32)	49.14
INT8 classification head only	50.57
INT8 regression head only	16.41
INT8 heads only (cls+reg)	17.09
INT8 all-except regression head	49.50
FP8 matched, 128-img calibration	48.50
FP32 reference (diagnostic build)	50.44

protocol validates the exclusion direction on both datasets, not the individual diagnostic values.

On the frozen paired protocol the selective arm validates the decomposition: INT8-selective matches FP8-matched within intervals that include zero on the clean control and on the 12-cell corrupt mean on KITTI (FP8 minus selective: J95 gap -0.28 [$-1.11, +0.56$], corrupt-mean gap $+0.17$ [$-0.07, +0.40$]; we note the J95 interval, -0.28 [$-1.11, +0.56$], spans over an AP point around zero, so the match is a bound rather than tight equivalence).

That match is a statement about the mean, not about every family: measured on the same ΔE basis (family gap minus the J95 gap), the selective residual widens significantly under motion blur ($+1.30$ [$+0.51, +2.04$]) and is null on JPEG (-0.11 [$-0.81, +0.50$]), so family-level divergence remains. On VOC—the cleaner check—a small residual remains on both bases: FP8-matched minus selective is $+0.42$ [$+0.13, +0.74$] on the J95 clean control and $+0.96$ [$+0.83, +1.09$] on the corrupt mean, widening under corruption (ΔE $+0.54$ [$+0.23, +0.82$]). The regression head is therefore the dominant contributor, but not the entire mechanism on VOC.

4.3. Corruption-aware calibration is itself context-dependent and does not repair the collapse

If the INT8-matched arm’s corruption fragility were driven by clean-only calibration ranges, the corrupt-calibration intervention should recover held-out corruption AP. Against the frozen criterion, the result splits by dataset (Table 5, Fig. 4). On VOC the held-out difference is $+0.02$ [$-0.13, +0.16$] AP, null, while the in-family cells seen in calibration improve by $+1.27$ [$+1.12, +1.42$]. On KITTI the primary criterion is met—held-out $+1.07$ [$+0.87, +1.26$]—but the gain closes less than a tenth of the held-out residual to FP8-matched on the same cells (INT8-corruptcalib minus FP8-matched: -13.34 [$-14.13, -12.57$]). Excluding the saturated motion-blur-s5 cell leaves both verdicts unchanged (held-out $+1.27$ [$+1.06, +1.50$] on KITTI; $+0.07$ [$-0.10, +0.23$] on VOC).

The paper’s own estimand explains why the KITTI gain is not demonstrably corruption-specific. Corrupt calibration also raises *clean-control* AP—by $+1.56$ [$+1.02, +2.10$] on KITTI and $+0.73$ [$+0.40, +1.01$] on VOC—so under the additive parallel-shift assumption of Section 3.6, part of the corrupted-cell gains would appear as a uniform accuracy shift unrelated to corruption. Subtracting that component, the corruption-specific held-out effect $\Delta E_{\text{held-out}}$ is -0.50 [$-1.01, +0.03$] on KITTI and -0.71 [$-1.02, -0.39$] on VOC: null-to-negative on both datasets. The KITTI held-out gain is therefore not demonstrably larger than the gain on uncorrupted images. Only the calibrated families show a corruption-specific benefit, and only on VOC ($\Delta E_{\text{in-family}}$ $+0.54$ [$+0.24, +0.88$]; KITTI $+0.05$ [$-0.54, +0.59$]). Calibration on corruptions fits the calibrated distribution and shifts overall accuracy slightly; in this fold it does not transfer robustness to unseen corruption families—whether that verdict generalizes is what the complementary fold tests next.

The codec-matched control arms settle the attribution question. Corrupt-calibration images carry a JPEG quality 95 terminal encoding that the clean-calibration images lack, so some of the gain could be codec adaptation. Calibrating the same 512 images at quality 95 without corruption (INT8-q95calib) produces no significant clean-AP change on either dataset (J95 clean -0.34 [$-0.93, +0.35$]

Recipe determinants of quantization fragility under corruption




Figure 3: RetinaNet recipe arms on KITTI and VOC: clean AP versus 12-cell corruption mean. INT8-matched collapses on KITTI (dataset-dependent); selective precision (INT8-sel = INT8-selective, reg-head in floating point) restores FP8-level behavior on KITTI and leaves a small residual on VOC; the calibration variants (INT8-q95 = codec-only calibration; INT8-cc = fold-1 corruption-aware calibration on noise and JPEG; INT8-cc2 = fold-2 calibration on fog and motion blur; a “sel-” prefix marks the selective recipe with the regression head in floating point) shift the matched arm little.

on KITTI, $+0.20$ $[-0.07, +0.48]$ on VOC) and does not improve the corrupt mean ($+0.04$ $[-0.10, +0.22]$ on KITTI; -0.17 $[-0.26, -0.09]$ on VOC, a small significant reduction). Against this control, the corrupt-calibration contrast is essentially unchanged in level terms (INT8-corruptcalib minus INT8-q95calib, held-out $+0.95$ $[+0.73, +1.15]$ on KITTI and $+0.13$ $[-0.02, +0.28]$ on VOC), and the corruption-specific held-out component is significantly negative on both datasets against this anchor ($\Delta E_{\text{held-out}}$ -0.94 $[-1.60, -0.33]$ on KITTI; -0.39 $[-0.69, -0.07]$ on VOC). On the selective tier, codec matching alone produces a small shift on corrupted cells that flips sign across datasets (corrupt mean $+0.36$ $[+0.23, +0.49]$ on KITTI; -0.12 $[-0.21, -0.04]$ on VOC) with no significant held-out corruption-specific component ($\Delta E_{\text{held-out}}$ $+0.05$ $[-0.48, +0.57]$ and -0.27 $[-0.56, +0.02]$), though the codec arm’s own in-family interaction is itself significantly negative on VOC (-0.30 $[-0.60, -0.02]$). Codec effects are therefore real but small in *level* terms—aggregate point estimates within half an AP point, far below the 21-AP collapse and the 1–2-AP intervention gains—and explain neither the full-INT8 collapse nor the level gains. On the interaction basis, however, codec effects are the same order as the corruption-specific components themselves (VOC q95-calib $\Delta E_{\text{held-out}}$ -0.31 $[-0.62, -0.02]$), and switching the anchor moves VOC fold-2’s $\Delta E_{\text{held-out}}$ from $+0.62$ $[+0.30, +0.96]$ to $+1.06$ $[+0.76, +1.38]$; the codec-anchored ΔE columns in Table 5 are therefore the relevant basis for the corruption-specific verdict, not an optional robustness check.

On the selective arm the pattern repeats at smaller magnitude. On VOC, corrupt calibration adds $+0.61$ $[+0.49, +0.72]$ AP on the corrupt mean, driven by the in-family cells ($+1.23$ $[+1.08, +1.39]$) with a null

held-out difference (-0.01 $[-0.17, +0.13]$), and improves the signed corrupt-mean residual to FP8-matched from -0.96 to -0.35 $[-0.46, -0.24]$ without closing it; the held-out residual is -0.51 $[-0.66, -0.38]$. The second frozen criterion—held-out residual to FP8-matched including zero—is therefore not met on either dataset (VOC -0.51 $[-0.66, -0.38]$; KITTI -1.30 $[-1.61, -1.00]$): calibration does not explain the residual. The fold-1 primary estimand also covers the selective-tier held-out contrast itself (sel-corruptcalib minus sel), which is -0.58 $[-0.80, -0.36]$ on KITTI and -0.01 $[-0.17, +0.13]$ on VOC—significantly negative on KITTI, under the contamination caveat of Section 3.4.

The complementary fold repeats the experiment with the family roles reversed—INT8-cc2calib calibrates on fog and motion blur and is tested on Gaussian noise and JPEG—and its result is not a second null: the intervention’s own effect is itself sign-unstable across the four fold $\times$ dataset cells (Table 5, Fig. 4). On KITTI the fold-2 arm is significantly *worse* on every level basis (J95 clean -2.71 $[-3.38, -1.93]$; calibrated fog+motion-blur cells -1.35 $[-1.56, -1.04]$; held-out -2.02 $[-2.31, -1.73]$), so the primary rule fails and the clean-level shift flips sign relative to fold 1. The degradation is largest on the clean control, so the corruption-specific components are positive on the calibrated families ($\Delta E_{\text{in-family}}$ $+1.36$ $[+0.61, +2.16]$) and null held-out ($+0.69$ $[-0.10, +1.37]$; codec-anchored $+0.37$ $[-0.28, +0.90]$): fold 2 damages this engine broadly rather than improving any basis. On VOC the same fold-2 recipe meets the primary rule: held-out $+1.66$ $[+1.52, +1.83]$ —recovering roughly half of the held-out residual to FP8-matched on those families (-1.74 $[-1.94, -1.53]$ remaining)—and, uniquely among the four cells, its corruption-specific component is positive and significant ($\Delta E_{\text{held-out}}$ $+0.62$ $[+0.30, +0.96]$, confirmed

against the codec anchor at +1.06 [+0.76, +1.38]). The selective tier does not repeat the split but diverges from it (all selective fold-2 contrasts are post-hoc; the frozen fold-2 rules cover only the full tier): on KITTI the fold-2 selective arm is significantly *lifted* where the full-INT8 arm is harmed (J95 +1.01 [+0.37, +1.72]; held-out +0.66 [+0.33, +0.92]; ΔE null; same development-contamination caveat as the fold-1 selective arm)—although neither lift is significant against its codec-matched baseline (sel-cc2 minus sel-q95, J95 +0.75 [−0.04, +1.56]; held-out +0.25 [−0.09, +0.52]), so codec adaptation can account for this particular lift. On VOC the selective held-out gain is +1.50 [+1.35, +1.66] and does survive its codec anchor (+1.64 [+1.49, +1.80]), with a corruption-specific component *larger* than the full tier’s ($\Delta E_{\text{held-out}}$ +0.86 [+0.55, +1.18], codec-anchored +1.16 [+0.87, +1.49])—against FP8-matched its held-out residual (INT8-sel-cc2calib minus FP8-matched) is +0.08 [−0.06, +0.23], the only cell in which the fold-1 analog of the second criterion (an interval including zero) is met; on KITTI fold 2 the same contrast excludes zero in the favorable direction (+1.04 [+0.64, +1.40], the selective fold-2 arm significantly exceeding FP8-matched on held-out, under the contamination caveat). Corruption-aware calibration therefore *can* transfer, but whether it does depends on the dataset×family partition: across the two folds its primary criterion is met on KITTI fold 1 and VOC fold 2 and fails on the other two cells, including a significant harm on KITTI fold 2.

The intervention verdict must be stated in two parts, separating the frozen rules from the interpretive analyses. Against the frozen criteria, the primary held-out test succeeds on KITTI fold 1 and VOC fold 2 and fails on the other two cells—including a significant *harm* on KITTI fold 2. Fold 1’s frozen secondary test (selective-versus-FP8 residual including zero) fails on both datasets; the fold-2 frozen corruption-specific rule (held-out ΔE excluding zero on the positive side against both anchors) is met on VOC and not met on KITTI; and the fold-2 selective-versus-FP8 contrast is a post-hoc analog, met on VOC and exceeded favorably on KITTI under the contamination caveat. It is worth stating plainly that on KITTI fold 1 the pre-specified reading of the primary rule is met, and the negative verdict rests on the post-hoc corruption-specific decomposition. Against that reading, fold 1 fails on both datasets while fold 2 produces a genuine positive component on VOC; codec adaptation cannot account for either direction of that instability. Two conclusions follow, and they should be kept apart. First, corruption-aware entropy calibration is itself context-dependent: it transfers a real corruption-specific benefit in one of four fold×dataset cells, leaves the others null or harmful, and never approaches the scale of the KITTI collapse—a second instance of the recipe×context dependence the paired layer showed. Second, corruption-aware calibration is not supported as a reliable repair for the corruption fragility, and clean-only calibration ranges are not supported as its dominant repairable cause: on KITTI, where the collapse lives, neither fold establishes a positive

held-out interaction—fold 1’s raw held-out gain closes under a tenth of the residual, and fold 2 degrades every level basis including the families it calibrated on (its $\Delta E_{\text{held-out}}$ point estimate is positive but null)—whereas excluding the regression head restores FP8-level behavior on KITTI and leaves only a small residual on VOC, without retraining.

A corruption-severity boundary is shared across precisions: at motion-blur severity 5 all matched-protocol arms collapse to below 6 AP on both datasets—KITTI: FP32 5.45, FP8-matched 5.59, INT8-matched 3.75, selective 4.52; VOC: 5.71/5.54/4.89/4.88 in the same order—FP32 included. That cell measures a common failure wall, not an INT8-specific deficit, and its equal weight in the corrupt mean compresses all contrasts built on it; the leave-one-out sensitivity removes it without changing any frozen-criterion verdict (two secondary codec-control contrasts do move across the zero boundary on VOC; Supplement S8).

4.4. Historical exploratory evidence

An earlier 144-cell exploratory grid analyzed the same paired YOLO11 engines—the same estimand ($G_c - G_{J95}$), codec control, pairing basis, and KITTI/VOC partitions—on four datasets including TT100K [66], differing in scope (the full COCO codec-control pool rather than the frozen 2,000-image subset, an extra dataset, and additional arms) and in the analysis plan (no Holm family, different summaries). It found a balanced additive interaction of −0.02 AP [−0.24, +0.15] across the grid and −0.55 [−0.78, −0.29] on post-selection VOC/KITTI holdouts evaluated with independently selected checkpoints, with 56 point-estimate sign reversals after clean adjustment—17 with intervals excluding zero on the negative side. These findings motivated the frozen paired analysis reported above but are not pooled with it. Notably, the historical post-selection holdout estimate is the one piece of this evidence built on checkpoints selected outside the evaluation partition, and its negative sign is consistent with several of the paired blocks. Two exploratory-era transfer cases—RT-DETR-L [67, 68] and legacy RetinaNet—exhibited catastrophic INT8 clean failures under the unmigrated recipe, evidence in retrospect of recipe fragility rather than format fragility.

5. Discussion

5.1. Recipe and context, not the format label

The three evidence layers converge on one conclusion: corruption robustness of a quantized detector is a property of the recipe, the model region, and the evaluation context, jointly with the format. No single layer could show this alone: the frozen paired layer shows the interaction has no stable sign (so pooling it into a format-level constant is meaningless); the diagnostic layer localizes *where* the fragility lives (the regression head); and the intervention layer shows that the *obvious remedy* (calibrating on corruptions) is itself context-dependent, transferring in one fold×dataset cell and harming another. Recipe choices—preprocessing contract, quantization sites—move

Table 5

RetinaNet intervention and codec-control contrasts (A minus B, AP points; plug-in point with 95% paired-bootstrap interval, 2,000 common-image resamples). J95 clean: codec-control clean cell; corrupt mean: 12 cells; in-family: Gaussian noise and JPEG (the fold-1 calibrated families); held-out: fog and motion blur (never calibrated in fold 1); for the q95-calibration rows, which see clean images only, these labels mark the evaluation partition rather than exposure; $\Delta E_{\text{held-out}}$ = held-out difference minus J95 clean difference, and $\Delta E_{\text{in-family}}$ is defined analogously. Fold-2 rows swap the family roles: their in-family is fog and motion blur and their held-out is Gaussian noise and JPEG. * interval excludes zero.

Contrast (A – B)	Dataset	J95 clean	Corrupt mean	In-family	Held-out	$\Delta E_{\text{held-out}}$
INT8-corruptcalib – INT8-matched	KITTI	+1.56* [+1.02, +2.10]	+1.34* [+1.17, +1.48]	+1.61* [+1.37, +1.79]	+1.07* [+0.87, +1.26]	–0.50 [–1.01, +0.03]
INT8-corruptcalib – INT8-matched	VOC	+0.73* [+0.40, +1.01]	+0.64* [+0.52, +0.76]	+1.27* [+1.12, +1.42]	+0.02 [–0.13, +0.16]	–0.71* [–1.02, –0.39]
INT8-q95calib – INT8-matched	KITTI	–0.34 [–0.93, +0.35]	+0.04 [–0.10, +0.22]	–0.02 [–0.23, +0.21]	+0.11 [–0.05, +0.33]	+0.45 [–0.24, +1.04]
INT8-q95calib – INT8-matched	VOC	+0.20 [–0.07, +0.48]	–0.17* [–0.26, –0.09]	–0.23* [–0.37, –0.11]	–0.11* [–0.228, –0.002]	–0.31* [–0.62, –0.02]
INT8-corruptcalib – INT8-q95calib	KITTI	+1.90* [+1.28, +2.54]	+1.29* [+1.08, +1.45]	+1.64* [+1.33, +1.86]	+0.95* [+0.73, +1.15]	–0.94* [–1.60, –0.33]
INT8-corruptcalib – INT8-q95calib	VOC	+0.52* [+0.22, +0.80]	+0.81* [+0.70, +0.93]	+1.50* [+1.35, +1.66]	+0.13 [–0.02, +0.28]	–0.39* [–0.69, –0.07]
INT8-sel-corruptcalib – INT8-selective	KITTI	–0.50 [–0.96, +0.03]	–0.50* [–0.65, –0.36]	–0.41* [–0.63, –0.24]	–0.58* [–0.80, –0.36]	–0.08 [–0.68, +0.41]
INT8-sel-corruptcalib – INT8-selective	VOC	+0.63* [+0.33, +0.94]	+0.61* [+0.49, +0.72]	+1.23* [+1.08, +1.39]	–0.01 [–0.17, +0.13]	–0.65* [–0.97, –0.35]
INT8-sel-q95calib – INT8-selective	KITTI	+0.26 [–0.15, +0.74]	+0.36* [+0.23, +0.49]	+0.41* [+0.21, +0.61]	+0.31* [+0.13, +0.54]	+0.05 [–0.48, +0.57]
INT8-sel-q95calib – INT8-selective	VOC	+0.16 [–0.10, +0.43]	–0.12* [–0.21, –0.04]	–0.14* [–0.27, –0.01]	–0.11 [–0.22, +0.01]	–0.27 [–0.56, +0.02]
INT8-sel-corruptcalib – INT8-sel-q95calib	KITTI	–0.76* [–1.35, –0.17]	–0.85* [–1.05, –0.69]	–0.82* [–1.08, –0.59]	–0.89* [–1.16, –0.65]	–0.13 [–0.83, +0.51]
INT8-sel-corruptcalib – INT8-sel-q95calib	VOC	+0.47* [+0.16, +0.76]	+0.73* [+0.61, +0.85]	+1.37* [+1.22, +1.54]	+0.09 [–0.07, +0.25]	–0.38* [–0.68, –0.08]
INT8-sel-corruptcalib – FP8-matched	KITTI	–0.23 [–1.12, +0.86]	–0.67* [–0.91, –0.44]	–0.03 [–0.37, +0.31]	–1.30* [–1.61, –1.00]	–1.08* [–2.18, –0.16]
INT8-sel-corruptcalib – FP8-matched	VOC	+0.21 [–0.06, +0.49]	–0.35* [–0.46, –0.24]	–0.19* [–0.32, –0.05]	–0.51* [–0.66, –0.38]	–0.72* [–1.00, –0.44]
INT8-corruptcalib – FP8-matched	KITTI	–19.51* [–20.90, –18.04]	–15.04* [–15.96, –14.17]	–16.74* [–17.85, –15.66]	–13.34* [–14.13, –12.57]	+6.17* [+5.12, +7.10]
INT8-corruptcalib – FP8-matched	VOC	–2.31* [–2.68, –1.98]	–2.13* [–2.30, –1.97]	–2.13* [–2.33, –1.93]	–2.12* [–2.32, –1.96]	+0.18 [–0.12, +0.50]
INT8-cc2calib – INT8-matched (fold 2)	KITTI	–2.71* [–3.38, –1.93]	–1.69* [–1.88, –1.44]	–1.35* [–1.56, –1.04]	–2.02* [–2.31, –1.73]	+0.69 [–0.10, +1.37]
INT8-cc2calib – INT8-matched (fold 2)	VOC	+1.04* [+0.72, +1.34]	+1.26* [+1.15, +1.37]	+0.86* [+0.72, +0.99]	+1.66* [+1.52, +1.83]	+0.62* [+0.30, +0.96]
INT8-cc2calib – INT8-q95calib (fold 2)	KITTI	–2.37* [–2.90, –1.74]	–1.73* [–1.96, –1.47]	–1.47* [–1.68, –1.18]	–2.00* [–2.32, –1.69]	+0.37 [–0.28, +0.90]
INT8-cc2calib – INT8-q95calib (fold 2)	VOC	+0.84* [+0.54, +1.12]	+1.43* [+1.32, +1.55]	+0.97* [+0.82, +1.12]	+1.89* [+1.75, +2.06]	+1.06* [+0.76, +1.38]
INT8-sel-cc2calib – INT8-selective (fold 2)	KITTI	+1.01* [+0.37, +1.72]	+0.62* [+0.42, +0.83]	+0.59* [+0.39, +0.86]	+0.66* [+0.33, +0.92]	–0.35 [–1.09, +0.28]
INT8-sel-cc2calib – INT8-selective (fold 2)	VOC	+0.63* [+0.34, +0.94]	+1.06* [+0.96, +1.18]	+0.63* [+0.48, +0.76]	+1.50* [+1.35, +1.66]	+0.86* [+0.55, +1.18]
INT8-sel-cc2calib – INT8-sel-q95calib (fold 2)	KITTI	+0.75 [–0.04, +1.56]	+0.27* [+0.08, +0.43]	+0.28* [+0.03, +0.52]	+0.25 [–0.09, +0.52]	–0.50 [–1.27, +0.33]
INT8-sel-cc2calib – INT8-sel-q95calib (fold 2)	VOC	+0.48* [+0.18, +0.75]	+1.19* [+1.07, +1.30]	+0.74* [+0.59, +0.87]	+1.64* [+1.49, +1.80]	+1.16* [+0.87, +1.49]
INT8-sel-cc2calib – FP8-matched (fold 2)	KITTI	+1.29* [+0.64, +2.03]	+0.45* [+0.20, +0.71]	–0.13 [–0.44, +0.20]	+1.04* [+0.64, +1.40]	–0.25 [–0.93, +0.36]
INT8-sel-cc2calib – FP8-matched (fold 2)	VOC	+0.22 [–0.06, +0.48]	+0.11 [–0.01, +0.22]	+0.13 [–0.02, +0.26]	+0.08 [–0.06, +0.23]	–0.14 [–0.40, +0.15]
INT8-cc2calib – FP8-matched (fold 2)	KITTI	–23.78* [–25.30, –22.13]	–18.07* [–19.07, –17.00]	–15.76* [–16.63, –14.86]	–20.37* [–21.58, –19.14]	+3.41* [+2.51, +4.23]
INT8-cc2calib – FP8-matched (fold 2)	VOC	–1.99* [–2.34, –1.66]	–1.51* [–1.68, –1.35]	–1.28* [–1.47, –1.12]	–1.74* [–1.94, –1.53]	+0.25 [–0.04, +0.57]




Figure 4: Calibration variants versus their clean-calibrated counterparts (AP points, 95% paired-bootstrap intervals). Grey: J95 clean control; blue: in-family corruptions; red: held-out corruptions. The family roles swap between folds: fold 1 calibrates Gaussian noise and JPEG and holds out fog and motion blur, while the groups marked “f2” (fold 2) calibrate fog and motion blur and hold out Gaussian noise and JPEG; for the q95-calibration rows, which see no corrupted images at all, the in-family/held-out labels only mark the evaluation partition. On KITTI fold 1 and the KITTI fold-2 full-INT8 tier the clean and held-out bars move together—the fold-2 full-tier movement being a harm, not a gain—so those held-out changes are not corruption-specific (the KITTI fold-2 selective tier instead improves significantly on both bases); VOC fold 1 moves the in-family and clean bars but leaves held-out null, and the VOC fold-2 groups are the exceptions, where held-out clearly exceeds clean for both the full and the selective tier. Codec-only calibration (q95-calib) shifts the full-INT8 arm’s aggregate bases by less than half an AP point; at per-cell granularity it moves individual cells significantly in both directions (a gain on the saturated KITTI motion-blur-s5 wall, a loss on a VOC noise cell, and several gains on the KITTI selective tier), so codec adaptation is real even where its aggregate contribution is small—a codec effect on shared failure cells rather than evidence of corruption coverage.

clean AP by 21–42 AP points, several-fold larger than the largest corruption-induced ΔE between FP8-matched and INT8-matched at the 12-cell corrupt-mean granularity (RetinaNet/KITTI: a -4.7 -AP narrowing). Interactions between arms that additionally differ in recipe run higher still—up to $+8.02$ AP on the in-family basis (INT8-cc2calib versus FP8-matched, KITTI)—and at per-family granularity the FP8–INT8-matched interaction itself reaches -11.4 AP (KITTI motion blur), so the “several-fold” comparison holds only at corrupt-mean granularity. Reporting “INT8” versus “FP8” robustness without the recipe contract is therefore not merely imprecise; it is capable of reversing conclusions.

5.2. The regression head is the fragile component—and why

The decomposition gives a clean localization (diagnostic-level, confirmed by the paired-protocol selective arm): classification-head quantization is harmless, regression-head quantization reproduces the collapse, and head exclusion restores clean-competitive accuracy. This is consistent with detector PTQ work emphasizing regression structures [27]. A plausible account, which our data are consistent with but do not isolate, is that box-regression outputs are unbounded deltas decoded through an exponential and a multiplication by the anchor size, so a fixed additive quantization error becomes a proportional localization error that AP at high IoU thresholds penalizes heavily, whereas classification scores pass through a saturating sigmoid and are thresholded. The intervention adds the sharper negative: on KITTI, where the collapse lives, corruption-inclusive calibration does not repair the head and produces no corruption-specific held-out benefit in either fold—the second fold harms every level basis—so the failure is consistent with a region-level quantization-error sensitivity that none of the tested calibration distributions repairs—although our design varies only the calibration image distribution, so head-local range-management choices beyond entropy calibration remain untested. The small corruption-specific transfer observed on VOC fold 2 shows calibration coverage is a real but unstable contributor, which the context-dependent framing already predicts. FP8, which quantizes the same head with an exponent field and therefore a wide dynamic range, holds 49.19 AP (the pre-study FP8-legacy engine) and 49.48 AP (FP8-matched) on KITTI clean inputs, which also fits a dynamic-range account. Separating weight from activation quantization within the head is the natural next ablation. Why the head-exclusion effect is about $8\times$ larger on KITTI than on VOC (21.3 versus 2.6 AP on the J95 basis) remains open; dataset-dependent target statistics (box-size distribution, class prior, average object scale) are the likely moderators.

5.3. Practical guidance

For deployment recipes, the evidence supports a concrete checklist: (i) verify the calibration preprocessing contract byte-for-byte against inference before interpreting any

quantization gap; (ii) under matched contracts, evaluate the regression head and classification head as separate quantization regions, in the spirit of sensitivity-based mixed-precision assignment [69]—head exclusion recovered > 20 AP here while promoting only the `/head/regression_head/` subgraph to floating point; (iii) do not rely on corruption-augmented calibration to transfer: across two folds its held-out effect ranged from a significant harm to a $+1.66$ -AP gain depending on the dataset $\times$ family partition, and codec-match the calibration input encoding to the deployment stream before attributing gains to corruption coverage; (iv) report paired interactions together with absolute corrupted AP and the clean gap, so floor effects cannot masquerade as robustness. Selective precision is a retraining-free default worth evaluating first: it matched FP8-matched within intervals on KITTI and trailed it by $+0.42$ J95-clean / $+0.96$ corrupt AP on VOC. On deployment cost we make no quantitative claim: wall-clock latency measured on our shared GPU (graphics processing unit) host is inadmissible as evidence under the frozen protocol, and engine-build determinism was not audited—the cost argument for head exclusion is that the promoted subgraph is confined to the regression-head nodes—a layerwise MAC count of the torchvision graph (`analysis/regression_head_macs.py`) puts it at 24.5–25.1% of convolution MACs at the frozen input size for the two reported class counts—while all remaining quantization sites stay INT8-configured in the recipe (kernel-level execution rates were not measured); an INT8-body/half-precision-head variant is untested here.

5.4. Limitations

The evidence is conditional on recorded artifacts: one checkpoint per dataset–family cell, one engine build per arm, one calibration method (entropy) and toolchain (ModelOpt/TensorRT) on the paired and intervention arms—the diagnostic sweeps additionally vary the estimator and the 128/512-image budget—and one runtime environment (TensorRT 11.1, RTX 5090); other PTQ methods may place quantization error differently. Bootstrap intervals cover image resampling only—not training-seed, calibration-draw, engine-build, or corruption-realization variance—and treat images as independent; KITTI frames drawn from shared driving sequences are correlated, so intervals there may be optimistic, and a sequence-clustered bootstrap is not computable from the retained manifests because drive metadata is not preserved in the evaluation ledger. Three artifact-provenance caveats bound the paired layer: the KITTI and VOC evaluation partitions coincide with the checkpoint-selection partitions for both detector families (and the COCO YOLO11 models use the official pretrained checkpoint), so absolute AP levels are optimistic; both paired arms share the selected checkpoint, but whether the two arms’ selection biases differ—and hence how much enters the gap and the interaction—is unquantified; the YOLO11 engines predate the frozen contract and carry TensorRT’s default TF32 state, while the RetinaNet quantization arms were built with TF32 disabled (the RetinaNet FP32 reference and legacy arms are

TF32-on `trtexec` builds, but the ΔE contrasts only join quantized arms built under the same policy), so the two families also differ in builder policy; and the paired protocol’s freeze covers the analysis contract (estimands, conditions, bootstrap schedule), not the artifacts it runs on—the nine-block Holm adjustment is an analysis-time choice implemented in the reporting code rather than a line in the frozen configuration. Holm adjustment covers the nine paired YOLO11 blocks; RetinaNet and intervention intervals are per contrast. Two-sided bootstrap p -values have resolution 1/1000. The corruption suite is synthetic and four types wide; severities order settings within a type and are not comparable across types, and the corrupt mean’s equal weighting includes the saturated motion-blur-s5 wall cell (reported alongside a leave-one-out sensitivity). The intervention uses two complementary folds over four synthetic corruption types; other partitions and non-synthetic shifts are untested, and each calibration cell saw only 75–102 images. The intervention arms also differ on the clean control (the corrupt-calibration gains include a significant uniform clean-AP shift), which is why the corruption-specific reading relies on the held-out ΔE decomposition rather than the raw held-out gain. The first fold’s frozen criteria were a held-out significance test and a selective-versus-FP8 residual test; the magnitude comparison and the held-out ΔE analysis were specified afterwards and are reported as interpretive. The second fold was frozen after that analysis and pre-specifies both rules. The codec controls attached to the fold-1 contrasts are post-hoc additions. Three further intervention caveats apply: the fold partition confounds corruption family with spectral character (fold 1 calibrates the high-frequency families and holds out the low-frequency ones, and fold 2 swaps them), “held-out JPEG” is not fully unseen in codec terms because every calibration image is already JPEG quality-95 encoded, and the many intervention contrasts carry no familywise multiplicity correction, so individual significant cells should be read against the full 28-row table rather than in isolation. The null-seeking selective-versus-FP8 criterion bounds the residual but specifies no equivalence margin, so an interval containing zero is not a proof of equivalence [70]. On the RetinaNet arms the J95 clean control departs from original clean by up to 0.8 AP (the fold-2 selective arm), larger than the ~ 0.3 -point fidelity seen on YOLO11. KITTI’s selective-arm numbers are development-contaminated; VOC serves as the cleaner check, and the two-dataset scope of the RetinaNet localization limits generality. The RetinaNet analysis covers one anchor-based architecture; RT-DETR engines exist only under the legacy recipe and were excluded from the paired protocol. The decomposition arms are diagnostic single points without intervals. Historical exploratory results share the paired YOLO11 engines; only their retrained-holdout and legacy-transfer arms use different builds, and all are reported only as context. COCO evidence is restricted to the frozen 2,000-image subset. Finally, AP is a single utility summary; downstream risk thresholds belong to the application.

6. Conclusion

A paired, codec-controlled protocol with a common-image bootstrap turns the question “is INT8 or FP8 more robust to corruption” into a question with a defensible answer: on the corruption-induced change in the gap there is no stable winner—the interaction changes sign across datasets and scales, while FP8 exceeds INT8 on absolute clean and corrupted AP in every matched block mean, and recipe and model region determine where the fragility lives. The RetinaNet case makes the localization concrete. A preprocessing contract fix recovers 21 AP points on KITTI; a regression-head exclusion recovers a further 21 on the same dataset—a magnitude that carries the disclosed development-contamination caveat, with its direction replicated on VOC; and a two-fold corruption-aware calibration intervention, read through the same paired estimand, is itself context-dependent—transferring a small corruption-specific benefit in one fold $\times$ dataset cell and significantly harming another—so calibration coverage is a secondary, unstable contributor rather than a reliable repair. Quantization robustness under corruption should be evaluated and reported per recipe—with paired inputs, clean-gap adjustment, and absolute corrupted AP—rather than per format label.

Data and code availability

The Supplementary Information restates the full frozen contract: estimands, condition schedule, arm registry, inference/evaluation constants, corruption-aware calibration schedule, bootstrap schedule, evidence layers, and a reporting checklist. Manuscript tables and figures are generated deterministically from frozen ledgers and manifests. Engine registries, calibration manifests (including the corrupt-calibration schedule), input manifests, metric records, and paired-bootstrap draw caches are hash-bound and retained in the accompanying compact-evidence package, from which the analysis source regenerates all reported numbers. The package is deposited in the versioned research archive under concept DOI <https://doi.org/10.5281/zenodo.22031663>, which currently resolves to the prior CVIU-era release lineage and mints the Neural Networks release on publication. Redistribution covers manifests, metrics, and analysis code; raw datasets, trained checkpoints, raw predictions, and built TensorRT engines are not redistributed.

Declaration of competing interest

The authors declare no competing interests.

CRedit authorship contribution statement

Dinh Thuan Nguyen: Conceptualization, Methodology, Software, Investigation, Formal analysis, Writing—original draft. **Lam Phuong Nguyen:** Validation, Data curation, Writing—review and editing. **Vinh Huy Nguyen:** Methodology, Visualization, Writing—review and editing. **Sy Vu**

Quang: Validation, Writing–review and editing. **Mohan Rajesh Elara:** Supervision, Writing–review and editing. **Anh Vu Le:** Supervision, Project administration, Writing–review and editing.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During preparation of this work, the authors used Cognition Devin for analysis-code development, experiment orchestration, and manuscript drafting support; OpenAI Codex for code refactoring, validation support, consistency checks, and language editing; OpenAI ChatGPT for literature-discovery support and manuscript critique; and Anthropic Claude Code for independent manuscript critique and submission-file review. These systems were not listed as authors or contributors and did not autonomously train detectors, build engines, generate predictions, fabricate or alter research data, or make final scientific decisions. The authors verified cited sources and numerical claims, reviewed every retained change, reran relevant validation scripts, and take full responsibility for the article.

Acknowledgments

This research received no specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

References

- [1] R. Krishnamoorthi, Quantizing deep convolutional networks for efficient inference: A whitepaper, arXiv preprint arXiv:1806.08342 (2018). doi:10.48550/arXiv.1806.08342.
- [2] M. Nagel, M. Fournarakis, R. A. Amjad, Y. Bondarenko, M. van Baalen, T. Blankevoort, A white paper on neural network quantization, arXiv preprint arXiv:2106.08295 (2021). doi:10.48550/arXiv.2106.08295.
- [3] A. Gholami, S. Kim, Z. Dong, Z. Yao, M. W. Mahoney, K. Keutzer, A survey of quantization methods for efficient neural network inference, arXiv preprint arXiv:2103.13630 (2021). doi:10.48550/arXiv.2103.13630.
- [4] B. Jacob, S. Kligys, B. Chen, M. Zhu, M. Tang, A. Howard, H. Adam, D. Kalenichenko, Quantization and training of neural networks for efficient integer-arithmetic-only inference, in: Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2018, pp. 2704–2713. doi:10.1109/CVPR.2018.00286.
- [5] R. Banner, Y. Nahshan, D. Soudry, Post training 4-bit quantization of convolutional networks for rapid-deployment, in: Advances in Neural Information Processing Systems, Vol. 32, 2019, pp. 7948–7956.
- [6] Z. Yao, Z. Dong, Z. Zheng, A. Gholami, J. Yu, E. Tan, L. Wang, Q. Huang, Y. Wang, M. W. Mahoney, K. Keutzer, HAWQ-V3: Dyadic neural network quantization, in: Proceedings of the 38th International Conference on Machine Learning, Vol. 139 of Proceedings of Machine Learning Research, 2021, pp. 11875–11886.
- [7] H. Wu, P. Judd, X. Zhang, M. Isaev, P. Micikevicius, Integer quantization for deep learning inference: Principles and empirical evaluation, arXiv preprint arXiv:2004.09602 (2020). doi:10.48550/arXiv.2004.09602.
- [8] P. Micikevicius, D. Stolic, N. Burgess, M. Cornea, P. Dubey, R. Grisenthwaite, S. Ha, A. Heinecke, P. Judd, J. Kamalu, N. Mellempudi, S. Oberman, M. Shoeybi, M. Siu, H. Wu, FP8 formats for deep learning, arXiv preprint arXiv:2209.05433 (2022). doi:10.48550/arXiv.2209.05433.
- [9] H. Shen, N. Mellempudi, X. He, Q. Gao, C. Wang, M. Wang, Efficient post-training quantization with FP8 formats, in: Proceedings of Machine Learning and Systems, Vol. 6, 2024, pp. 483–498.
- [10] D. Hendrycks, T. Dietterich, Benchmarking neural network robustness to common corruptions and perturbations, in: International Conference on Learning Representations, 2019.
- [11] C. Michaelis, M. Mitzkus, R. Geirhos, E. Rusak, O. Bringmann, A. S. Ecker, M. Bethge, W. Brendel, Benchmarking robustness in object detection: Autonomous driving when winter is coming, arXiv preprint arXiv:1907.07484 (2019). doi:10.48550/arXiv.1907.07484.
- [12] C. Kamann, C. Rother, Benchmarking the robustness of semantic segmentation models, in: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2020, pp. 8828–8838.
- [13] J. Liu, Z. Wang, L. Ma, C. Fang, T. Bai, X. Zhang, J. Liu, Z. Chen, Benchmarking object detection robustness against real-world corruptions, International Journal of Computer Vision 132 (10) (2024) 4398–4416. doi:10.1007/s11263-024-02096-6.
- [14] M. van Baalen, A. Kuzmin, S. S. Nair, Y. Ren, E. Mahurin, C. Patel, S. Subramanian, S. Lee, M. Nagel, J. Soriaga, T. Blankevoort, FP8 versus INT8 for efficient deep learning inference, arXiv preprint arXiv:2303.17951 (2023). doi:10.48550/arXiv.2303.17951.
- [15] G. Jocher, J. Qiu, Ultralytics YOLO11, official software documentation and citation record; accessed 14 August 2026 (2024). URL <https://docs.ultralytics.com/models/yolo11/>
- [16] T.-Y. Lin, P. Goyal, R. Girshick, K. He, P. Dollár, Focal loss for dense object detection, in: Proceedings of the IEEE International Conference on Computer Vision, 2017, pp. 2980–2988. doi:10.1109/ICCV.2017.324.
- [17] PyTorch Core Team, Torchvision 0.28 RetinaNet ResNet-50 FPN V2 documentation, Official software documentation, version 0.28; accessed 22 August 2026 (2026). URL https://docs.pytorch.org/vision/0.28/models/generated/torchvision.models.detection.retinanet_resnet50_fpn_v2.html
- [18] M. Nagel, M. van Baalen, T. Blankevoort, M. Welling, Data-free quantization through weight equalization and bias correction, in: Proceedings of the IEEE/CVF International Conference on Computer Vision, 2019, pp. 1325–1334.
- [19] Y. Cai, Z. Yao, Z. Dong, A. Gholami, M. W. Mahoney, K. Keutzer, ZeroQ: A novel zero shot quantization framework, in: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2020, pp. 13169–13178.
- [20] M. Nagel, R. A. Amjad, M. van Baalen, C. Louizos, T. Blankevoort, Up or down? adaptive rounding for post-training quantization, in: Proceedings of the 37th International Conference on Machine Learning, Vol. 119 of Proceedings of Machine Learning Research, 2020, pp. 7197–7206.
- [21] I. Hubara, Y. Nahshan, Y. Hanani, R. Banner, D. Soudry, Accurate post training quantization with small calibration sets, in: Proceedings of the 38th International Conference on Machine Learning, Vol. 139 of Proceedings of Machine Learning Research, 2021, pp. 4466–4475.
- [22] Y. Li, R. Gong, X. Tan, Y. Yang, P. Hu, Q. Zhang, F. Yu, W. Wang, S. Gu, BRECQ: Pushing the limit of post-training quantization by block reconstruction, in: International Conference on Learning Representations, 2021.
- [23] X. Wei, R. Gong, Y. Li, X. Liu, F. Yu, QDrop: Randomly dropping quantization for extremely low-bit post-training quantization, in: International Conference on Learning Representations, 2022. URL <https://openreview.net/forum?id=ySQH0oDyp7>
- [24] K. Oksuz, B. C. Cam, S. Kalkan, E. Akbas, Imbalance problems in object detection: A review, IEEE Transactions on Pattern Analysis and Machine Intelligence 43 (10) (2021) 3388–3415. doi:10.1109/TPAMI.2020.2981890.
- [25] J. Liu, L. Niu, Z. Yuan, D. Yang, X. Wang, W. Liu, PD-Quant: Post-training quantization based on prediction difference metric, in: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2023, pp. 24427–24437. doi:10.1109/CVPR52729.

2023.02340.

- [26] L. Niu, J. Liu, Z. Yuan, D. Yang, X. Wang, W. Liu, Improving post-training quantization on object detection with task loss-guided L_p metric, arXiv preprint arXiv:2304.09785 (2023). doi:10.48550/arXiv.2304.09785.
- [27] Y. Ding, W. Feng, C. Chen, J. Guo, X. Liu, Reg-PTQ: Regression-specialized post-training quantization for fully quantized object detector, in: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2024, pp. 16174–16184.
- [28] M. Kim, D. Lee, J. Yu, J. Hur, G. Kim, J. Kim, Inlier-centric post-training quantization for object detection models, in: International Conference on Learning Representations, 2026. URL <https://openreview.net/forum?id=GN9otzf5o6>
- [29] S. K. Esser, J. L. McKinstry, D. Bablani, R. Appuswamy, D. S. Modha, Learned step size quantization, in: International Conference on Learning Representations, 2020. URL <https://openreview.net/forum?id=rkgO66VKDS>
- [30] Y. Wei, X. Pan, H. Qin, W. Ouyang, J. Yan, Quantization mimic: Towards very tiny CNN for object detection, in: Computer Vision–ECCV 2018, 2018, pp. 267–283.
- [31] R. Li, Y. Wang, F. Liang, H. Qin, J. Yan, Fully quantized network for object detection, in: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2019, pp. 2810–2819.
- [32] P. Chen, J. Liu, B. Zhuang, M. Tan, C. Shen, AQD: Towards accurate quantized object detection, in: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2021, pp. 104–113.
- [33] S. Nguyen-Tien, T. Le-Tien, H. Ly-Thanh, N. Ly-Tu, ONNX-based architectures for post-training quantization face detection on edge devices, Advances in Electrical and Electronic Engineering Early Access (2026). URL <https://advances.vsb.cz/early-access>
- [34] N. Wang, J. Choi, D. Brand, C.-Y. Chen, K. Gopalakrishnan, Training deep neural networks with 8-bit floating point numbers, in: Advances in Neural Information Processing Systems, Vol. 31, 2018.
- [35] X. Sun, J. Choi, C.-Y. Chen, N. Wang, S. Venkataramani, V. Srinivasan, X. Cui, W. Zhang, K. Gopalakrishnan, Hybrid 8-bit floating point (HFP8) training and inference for deep neural networks, in: Advances in Neural Information Processing Systems, Vol. 32, 2019.
- [36] B. Nouné, P. Jones, D. Justus, D. Masters, C. Luschi, 8-bit numerical formats for deep neural networks, arXiv preprint arXiv:2206.02915 (2022). doi:10.48550/arXiv.2206.02915.
- [37] A. Kuzmin, M. van Baalen, Y. Ren, M. Nagel, J. Peters, T. Blankevoort, FP8 quantization: The power of the exponent, in: Advances in Neural Information Processing Systems, Vol. 35, 2022. doi:10.52202/068431-1065.
- [38] Z. Zhang, Y. Zhang, G. Shi, Y. Shen, R. Gong, X. Xia, Q. Zhang, L. Lu, X. Liu, Exploring the potential of flexible 8-bit format: Design and algorithm, arXiv preprint arXiv:2310.13513 (2023). doi:10.48550/arXiv.2310.13513.
- [39] T. Karimov, H. Imani, A. Kazakov, Quantization robustness to input degradations for object detection, in: 2025 Innovations in Intelligent Systems and Applications Conference (ASYU), IEEE, 2025. doi:10.1109/ASYU67174.2025.11208310.
- [40] Q. Li, Y. Meng, C. Tang, J. Jiang, Z. Wang, Investigating the impact of quantization on adversarial robustness, arXiv preprint arXiv:2404.05639 (2024). doi:10.48550/arXiv.2404.05639.
- [41] J. Xiao, Z. Li, L. Yang, Y. Mei, Q. Gu, TTAQ: Towards stable post-training quantization in continuous domain adaptation, arXiv preprint arXiv:2412.09899 (2024). doi:10.48550/arXiv.2412.09899.
- [42] H. Yaghoubi Araghi, P. Pilevar, M. C. Lin, Recti-Q: Feature-space rectification for out-of-distribution-robust quantized perception in edge robotics, arXiv preprint arXiv:2607.18540 (2026). doi:10.48550/arXiv.2607.18540.
- [43] D. Wang, E. Shelhamer, S. Liu, B. Olshausen, T. Darrell, Tent: Fully test-time adaptation by entropy minimization, in: International Conference on Learning Representations, 2021. URL <https://openreview.net/forum?id=uX13bZLkr3c>
- [44] Q. Wang, O. Fink, L. Van Gool, D. Dai, Continual test-time domain adaptation, in: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2022, pp. 7201–7211.
- [45] S. Niu, J. Wu, Y. Zhang, Y. Chen, S. Zheng, P. Zhao, M. Tan, Efficient test-time model adaptation without forgetting, in: Proceedings of the 39th International Conference on Machine Learning, Vol. 162 of Proceedings of Machine Learning Research, 2022, pp. 16888–16905.
- [46] S. Schneider, E. Rusak, L. Eck, O. Bringmann, W. Brendel, M. Bethge, Improving robustness against common corruptions by covariate shift adaptation, in: Advances in Neural Information Processing Systems, Vol. 33, 2020, pp. 11539–11551.
- [47] D. Hendrycks, N. Mu, E. D. Cubuk, B. Zoph, J. Gilmer, B. Lakshminarayanan, AugMix: A simple data processing method to improve robustness and uncertainty, in: International Conference on Learning Representations, 2020.
- [48] E. Mintun, A. Kirillov, S. Xie, On interaction between augmentations and corruptions in natural corruption robustness, in: Advances in Neural Information Processing Systems, Vol. 34, 2021.
- [49] D. Hoiem, Y. Chodpathumwan, Q. Dai, Diagnosing error in object detectors, in: Computer Vision–ECCV 2012, 2012, pp. 340–353. doi:10.1007/978-3-642-33712-3_25.
- [50] D. Bolya, S. Foley, J. Hays, J. Hoffman, TIDE: A general toolbox for identifying object detection errors, in: Computer Vision–ECCV 2020, 2020, pp. 558–573. doi:10.1007/978-3-030-58580-8_33.
- [51] N. A. Thacker, A. F. Clark, J. L. Barron, J. R. Beveridge, P. Courtney, W. R. Crum, V. Ramesh, C. Clark, Performance characterization in computer vision: A guide to best practices, Computer Vision and Image Understanding 109 (3) (2008) 305–334. doi:10.1016/j.cviu.2007.04.006.
- [52] L. Wen, D. Du, Z. Cai, Z. Lei, M.-C. Chang, H. Qi, J. Lim, M.-H. Yang, S. Lyu, UA-DETRAC: A new benchmark and protocol for multi-object detection and tracking, Computer Vision and Image Understanding 193 (2020) 102907. doi:10.1016/j.cviu.2020.102907.
- [53] NVIDIA Corporation, NVIDIA TensorRT 11.1.0 release notes, Official software documentation, the recorded build reports version 11.1.0.106; accessed 3 September 2026 (2026). URL <https://docs.nvidia.com/deeplearning/tensorrt/latest/getting-started/release-notes-11/11.1.0.html>
- [54] NVIDIA Corporation, NVIDIA Model Optimizer 0.45.0 release, Official software release, version 0.45.0; accessed 22 August 2026 (2026). URL <https://github.com/NVIDIA/Model-Optimizer/releases/tag/0.45.0>
- [55] S. Migacz, 8-bit inference with TensorRT, NVIDIA GPU Technology Conference (2017).
- [56] J. D. Angrist, J.-S. Pischke, Mostly Harmless Econometrics: An Empiricist’s Companion, Princeton University Press, 2009.
- [57] A. Geiger, P. Lenz, R. Urtasun, Are we ready for autonomous driving? The KITTI vision benchmark suite, in: Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2012, pp. 3354–3361. doi:10.1109/CVPR.2012.6248074.
- [58] M. Everingham, S. M. A. Eslami, L. Van Gool, C. K. I. Williams, J. Winn, A. Zisserman, The PASCAL visual object classes challenge: A retrospective, International Journal of Computer Vision 111 (1) (2015) 98–136. doi:10.1007/s11263-014-0733-5.
- [59] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, C. L. Zitnick, Microsoft COCO: Common objects in context, in: Computer Vision–ECCV 2014, 2014, pp. 740–755. doi:10.1007/978-3-319-10602-1_48.
- [60] X. Bouthillier, P. Delaunay, M. Bronzi, A. Trofimov, B. Nichyporuk, J. Szeto, N. Mohammadi Sepahvand, E. Raff, K. Madan, V. Voleti, S. Ebrahimi Kahou, V. Michalski, T. Arbel, C. Pal, G. Varoquaux, P. Vincent, Accounting for variance in machine learning benchmarks, in: Proceedings of Machine Learning and Systems, Vol. 3, 2021.
- [61] J. Pineau, P. Vincent-Lamarre, K. Sinha, V. Larivière, A. Beygelzimer, F. d’Alché Buc, E. Fox, H. Larochelle, Improving reproducibility in machine learning research: A report from the NeurIPS 2019 reproducibility program, Journal of Machine Learning Research

- 22 (164) (2021) 1–20.
- [62] B. Efron, Bootstrap methods: Another look at the jackknife, *The Annals of Statistics* 7 (1) (1979) 1–26. doi:10.1214/aos/1176344552.
- [63] S. Holm, A simple sequentially rejective multiple test procedure, *Scandinavian Journal of Statistics* 6 (2) (1979) 65–70.
- [64] W. G. Cochran, The combination of estimates from different experiments, *Biometrics* 10 (1) (1954) 101–129. doi:10.2307/3001666.
- [65] J. P. T. Higgins, S. G. Thompson, Quantifying heterogeneity in a meta-analysis, *Statistics in Medicine* 21 (11) (2002) 1539–1558. doi:10.1002/sim.1186.
- [66] Z. Zhu, D. Liang, S. Zhang, X. Huang, B. Li, S. Hu, Traffic-sign detection and classification in the wild, in: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 2110–2118.
- [67] Y. Zhao, W. Lv, S. Xu, J. Wei, G. Wang, Q. Dang, Y. Liu, J. Chen, DETRs beat YOLOs on real-time object detection, in: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024, pp. 16965–16974.
- [68] N. Carion, F. Massa, G. Synnaeve, N. Usunier, A. Kirillov, S. Zagoruyko, End-to-end object detection with transformers, in: *Computer Vision–ECCV 2020*, 2020, pp. 213–229. doi:10.1007/978-3-030-58452-8_13.
- [69] Z. Dong, Z. Yao, A. Gholami, M. W. Mahoney, K. Keutzer, HAWQ: Hessian aware quantization of neural networks with mixed-precision, in: *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2019, pp. 293–302.
- [70] D. Lakens, Equivalence tests: A practical primer for t tests, correlations, and meta-analyses, *Social Psychological and Personality Science* 8 (4) (2017) 355–362. doi:10.1177/1948550617697177.